



Analytics Solutions Bootcamp: #07

Artificial Intelligence Essentials

Designing practical AI-powered analytics solutions

Hello, everyone! Welcome to the seventh and final module of the Analytics Solutions Bootcamp. Today, you will learn about the essentials of artificial intelligence or AI.

Topics to be covered



- What is artificial intelligence?
- Introduction to artificial neural networks
- Generative AI and prompt engineering
- Building AI applications: RAGs and Agents
- Evaluating AI solutions
- Responsible AI and AI governance


What is artificial intelligence?

Getting grounded with artificial intelligence concepts
and how we can use it

What is artificial intelligence?

Artificial intelligence in a nutshell >>

Artificial Intelligence (AI) is a form of technology that lets machines solve problems in a way that is identical to how people do it. It helps businesses make better decisions on a large scale by helping them recognize images, create content, and make predictions based on data.



September 11, 2026

Analytics Solutions Bootcamp

4

<https://www.freecodecamp.org/news/machine-learning-vs-deep-learning-vs-generative-ai/>

What is artificial intelligence?

AI use cases across industries >>

- **Banking & FinTech:** fraud detection, AI customer service assistants
- **Retail & E-Commerce:** personalized product recommendations, demand forecasting
- **Healthcare:** medical imaging analysis, clinical documentation assistants
- **Manufacturing:** predictive maintenance, automated visual quality inspection
- **Logistics & Supply Chain:** route optimization, inventory and demand prediction
- **Telecommunications:** customer churn prediction, AI-powered network monitoring

What is artificial intelligence?

AI use cases across industries >>

- **Human Resources:** employee support assistants, workforce analytics
- **Marketing:** personalized content generation, customer segmentation
- **Insurance:** claims processing automation, fraud and risk detection
- **Education:** AI tutors, personalized learning content
- **Real Estate:** property valuation, property/customer inquiry assistants
- **Government / Public Sector:** citizen-service assistants, document processing

What is artificial intelligence?

Three perspectives of AI solutions >>

Predictive AI

Predict an outcome

Will this customer
churn?

Generative AI

Generate new content

Summarize this
complaint.

Agentic AI

**Use information and
tools to accomplish
tasks**

Investigate the
complaint and
recommend the next
action.


Start with the business problem — not with the technology.

Introduction to artificial neural networks

A conceptual view of artificial neural networks and how it is used for different problems


Artificial neural networks (ANNs) >>

Artificial neural networks (ANNs) consist of interconnected artificial neurons organized into layers that learn from training data to recognize patterns, make predictions and solve complex problems.



ARTIFICIAL NEURAL NETWORK – HIGH LEVEL STRUCTURE

An ANN learns a function that maps inputs to outputs through interconnected neurons.




Explain the diagram

This makes the ANN a black box model, which means that we pass an input and get an output without knowing fully what happened behind the scenes. We can build it from scratch and explain what happens in each hidden layer, but it won't be intuitive for business use.

Introduction to artificial neural networks

Example: Will the customer buy? >>



September 11, 2026

Analytics Solutions Bootcamp

12

Activation functions introduce non-linearity to the model.

ReLU – Rectified Linear Unit

Note that this is just the first iteration of the neural network. In other words, we initialized the model with random weights and biases, hence, these weights and biases are not yet optimal for prediction.

How the network learns?



A neural network model learns by going through the following process.

1. Forward propagation: pass input to the network to calculate predictions.
2. Calculate loss: Using a loss function, calculate the error in model.
3. Backward propagation: calculates how much does each weight contribute to the error.
4. Weight update: adjusts the weights and biases to minimize error.

Calculating loss



Loss function tracks the degree of error in a model's outputs.

- Binary Cross-Entropy (Log Loss) is for binary classification.
 - $L = -(y \ln(\hat{y}) + (1 - y) \ln(1 - \hat{y}))$
- Specific functions for other prediction problems (regression, multiclass)

$$L = -(y \ln(\hat{y}) + (1 - y) \ln(1 - \hat{y}))$$

Where: $y = 1$ and $\hat{y} = 0.5167$

$$L = -(1 \ln(0.5167) + (1 - 1) \ln(1 - 0.5167))$$

$$L = -(\ln(0.5167) + (0) \ln(0.4833))$$

$$L = -(\ln(0.5167))$$

$$L = -(-0.6603)$$

$$L = \boxed{0.6603}$$

We have different loss functions. We select the loss function based on our prediction problem.

- For regression, we can use MSE, MAE, Huber Loss
- For classification:
 - For binary classification, we can use Binary Cross-Entropy or Log Loss
 - For multi-class classification, we can use Categorical Cross-Entropy

Backward propagation



- We calculate how much did each weight contribute to the error.
- Uses chain rule to differentiate the weights throughout the network.
- Calculate the output layer error signal > calculate gradients for output weights and biases >> calculate gradients for neuron weights and biases

Calculate output error signal



Calculate the output-layer error signal

$$\frac{\partial L}{\partial z_0} = \hat{y} - y$$

$$\frac{\partial L}{\partial z_0} = 0.5167 - 1$$

$$\frac{\partial L}{\partial z_0} = -0.4833$$

Get output weight gradients



Calculate gradient for output weight v_1

$$\begin{aligned}\frac{\partial L}{\partial v_1} &= \frac{\partial L}{\partial z_0} \times h_1 \\ \frac{\partial L}{\partial v_1} &= -0.4833 \times 0.63 \\ \frac{\partial L}{\partial v_1} &= -0.3045\end{aligned}$$

Calculate gradient for output weight v_2

$$\begin{aligned}\frac{\partial L}{\partial v_2} &= \frac{\partial L}{\partial z_0} \times h_2 \\ \frac{\partial L}{\partial v_2} &= -0.4833 \times 0 \\ \frac{\partial L}{\partial v_2} &= 0\end{aligned}$$

Get output bias gradients



Calculate gradient for output bias b_1

$$\frac{\partial L}{\partial v_1} = \frac{\partial L}{\partial z_0} \times \frac{\partial z_0}{\partial b_0}$$

$$\frac{\partial L}{\partial v_1} = (-0.4833)(1)$$

$$\frac{\partial L}{\partial v_1} = \boxed{-0.4833}$$

Get neuron 1 weight gradients >>

Send error back to neuron 1

$$\frac{\partial L}{\partial h_1} = \frac{\partial L}{\partial z_0} \times v_1$$
$$\frac{\partial L}{\partial h_1} = (-0.4833)(0.9) \approx -0.4350$$
$$\frac{\partial L}{\partial z_1} = \frac{\partial L}{\partial h_1} \times \text{ReLU}'(z_1)$$
$$\frac{\partial L}{\partial z_1} = (-0.4350)(1) \approx \boxed{-0.4350}$$

Get neuron 1 weight gradient

$$\frac{\partial L}{\partial w_{11}} = \frac{\partial L}{\partial z_1} \times x_1, \frac{\partial L}{\partial w_{12}} = \frac{\partial L}{\partial z_2} \times x_2, \frac{\partial L}{\partial w_{13}} = \frac{\partial L}{\partial z_3} \times x_3$$
$$\frac{\partial L}{\partial w_{11}} = (-0.4350)(0.8) \approx \boxed{-0.3479}$$
$$\frac{\partial L}{\partial w_{12}} = (-0.4350)(0.1) \approx \boxed{-0.0435}$$
$$\frac{\partial L}{\partial w_{13}} = (-0.4350)(0.4) \approx \boxed{-0.1740}$$

Get neuron 1 bias gradients



Calculate gradient for output bias b_1

$$\frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial z_0} \times \frac{\partial z_0}{\partial h_1} \times \frac{\partial h_1}{\partial z_1} \times \frac{\partial z_1}{\partial b_1}$$

$$\frac{\partial L}{\partial b_1} = (-0.4833)(0.9)(1)(1)$$

$$\frac{\partial L}{\partial b_1} = -0.4350$$

Get neuron 2 weight gradients >>

Send error back to neuron 1

$$\begin{aligned}\frac{\partial L}{\partial h_2} &= \frac{\partial L}{\partial z_0} \times v_2 \\ \frac{\partial L}{\partial h_2} &= (-0.4833)(0.7) \approx -0.3383 \\ \frac{\partial L}{\partial z_2} &= \frac{\partial L}{\partial h_2} \times \text{ReLU}'(z_2) \\ \frac{\partial L}{\partial z_2} &= (-0.3383)(0) \approx 0\end{aligned}$$

Get neuron 1 weight gradient

$$\begin{aligned}\frac{\partial L}{\partial w_{21}} &= \frac{\partial L}{\partial z_1} \times x_1, \frac{\partial L}{\partial w_{22}} = \frac{\partial L}{\partial z_2} \times x_2, \frac{\partial L}{\partial w_{23}} = \frac{\partial L}{\partial z_3} \times x_3 \\ \frac{\partial L}{\partial w_{21}} &= (0)(0.8) \approx 0 \\ \frac{\partial L}{\partial w_{22}} &= (0)(0.1) \approx 0 \\ \frac{\partial L}{\partial w_{23}} &= (0)(0.4) \approx 0\end{aligned}$$

Get neuron 1 bias gradients



Calculate gradient for output bias b_2

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial z_0} \times \frac{\partial z_0}{\partial h_2} \times \frac{\partial h_2}{\partial z_2} \times \frac{\partial z_2}{\partial b_2}$$

$$\frac{\partial L}{\partial b_1} = (-0.4833)(0.7)(0)(1)$$

$$\frac{\partial L}{\partial b_1} = 0$$

Backward propagation summary >>

Final gradients from backpropagation	
Weights	Gradient
Output (v_1, v_2)	-0.3045, 0
Neuron 1 weights (w_{11}, w_{12}, w_{13})	-0.3479, -0.0435, -0.1740
Neuron 2 weights (w_{21}, w_{22}, w_{23})	0, 0, 0
Neuron 1 bias	-0.4350
Neuron 2 bias	0
Output bias	-0.4833

Updating the weights



Now that backpropagation has calculated the gradients, **gradient descent** actually updates the weights and biases.

- Learning rate: $\eta = 0.1$
- Update rule: $w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$


Calculating new weights




Final gradients from backpropagation			
Parameter	Old value	Gradient	New value
Output v_1	0.9	-0.3045	0.93045
Neuron 1 w_1	0.7	-0.3479	0.73479
Neuron 1 w_2	0.5	-0.0435	0.50435
Neuron 1 w_3	-0.2	-0.1740	-0.18260
Neuron 1 bias	0.1	-0.4350	0.14350
Output bias	-0.5	-0.4833	-0.45167

Introduction to artificial neural networks

Example: Weight adjusted model >>




Predicting Titanic survival probability



Predicting Iris flower specie

Iris Dataset – Multi-Class Classification


Predict the Iris species (setosa, versicolor, virginica)




Estimating diamond prices

Diamonds Dataset – Regression

Predict the price (in USD) of a diamond




Sentiment analysis with ANNs



Q&A with artificial neural network

Sentence Prediction with Neural Networks

Example: Answer the question



From ANNs to generative AI >>

We can think of **generative AI (Gen AI)** models as large-scale neural networks designed to create various contents. A **large language model (LLM)** is a type of generative AI that can understand and write human language.

- Umbrella ANN model has 11 trainable parameters.
- An LLMs may have **billions of trainable parameters**.
 - Meta's Llama 3 has **70 billion parameters**

Generative AI and prompt engineering

Understanding how generative AI works and how we can send prompts for better outputs

September 11, 2026

Analytics Solutions Bootcamp

33

What is generative AI?



Generative AI (gen AI) is artificial intelligence (AI) that can create original content such as text, images, video, audio or software code in response to a user's prompt or request.

Predictive AI (machine learning or deep learning models) predicts based on limited set of values and outcomes. Generative AI generates new diverse outputs.

How generative AI works?



Training

Training on huge volumes of raw data through deep learning models to build the foundation “generalist” model

Tuning

Tuning the foundation model for specific use case by feeding new documents and performing reinforcement learning with human feedback

Generation, evaluation, retuning

Gen AI model is now ready for generation. Models constantly learn from user feedback through actual conversations.

Large language models (LLMs)



Large language models (LLMs) are trained on immense amounts of data, making them capable of understanding and generating natural language and other types of content to perform a wide range of tasks. LLMs are built on a **transformer** which excels at handling sequences of words and capturing patterns in text.






Image generation models



Diffusion models are generative models used primarily for image generation and other computer vision tasks. Diffusion-based neural networks are trained through deep learning to progressively “diffuse” samples with random noise, then reverse that diffusion process to generate high-quality images.



Generating an image of a dog



Multimodal AI models



Multimodal AI refers to machine learning models capable of processing and integrating information from multiple modalities or types of data. These modalities can include text, images, audio, video and other forms of sensory input.

Prompt engineering



A **prompt** is a cue, instruction, or question that guides a person or an AI tool to take action or produce a response. A prompt can contain information like the **instruction** or **question** you are passing to the model and include other details such as **context**, **inputs**, or **examples**. You can use these elements to instruct the model more effectively to improve the quality of results.


Introduction to generative AI

Prompt engineering examples



Tell me about artificial intelligence

Artificial intelligence (AI) is the broad field of creating computer systems that can perform tasks that normally require human intelligence—such as understanding language, recognizing patterns, making predictions, solving problems, and generating content.



A useful way to think about AI is as an umbrella term. Several important technologies sit underneath it:

- **Artificial Intelligence (AI)** — the broadest category: machines performing intelligent tasks.
- **Machine Learning (ML)** — a subset of AI where systems learn patterns from data rather than relying only on explicitly programmed rules.
- **Deep Learning (DL)** — a subset of ML that uses artificial neural networks with many layers to learn complex patterns.
- **Large Language Models (LLMs)** — deep-learning models trained on enormous amounts of text and other data to understand and generate language.
- **Generative AI** — AI designed to create new content, such as text, images, audio, video, or code.

Explain artificial intelligence to a 12-year-old in 3 short paragraphs. Use one real-world analogy and give 3 everyday examples. Avoid technical jargon.

Artificial intelligence, or **AI**, is a way of making computers do things that normally need human thinking. AI can recognize pictures, understand questions, suggest things you might like, and even create stories or images. It learns by looking at lots of examples and finding patterns.

Think of AI like a **student learning to recognize animals**. If you show the student hundreds of pictures of cats and dogs, they eventually learn what makes a cat look different from a dog. AI learns in a similar way: it studies many examples and uses what it has learned to make a good guess when it sees something new.

You probably already use AI every day. YouTube or Netflix recommends videos and movies you might enjoy, **your phone can recognize your face to unlock itself**, and Google Maps can suggest a **faster route** when there is traffic. AI isn't a person and doesn't think exactly like we do, but it can be very useful for helping people solve problems and complete tasks.

🗨️ 🔍 🔄 ⋮

September 11, 2026

Analytics Solutions Bootcamp

41

Types of prompts



- **Direct prompts (Zero-shot):** instructing or asking the model without any context or examples
- **One-, few- and multi-shot prompts:** provides the model with one or more examples of the desired input-output pairs before presenting the actual prompt
- **Chain of Thought Prompts:** CoT prompting encourages the model to break down complex reasoning into a series of intermediate steps, leading to a more comprehensive and well-structured final output.

Chain of thought prompt example >>

A store sold 120 units in January and 150 units in February. Calculate the percentage increase. Show the calculation step by step, then give the final answer in one sentence.

Step 1: Find the increase in units.

$$150 - 120 = 30 \text{ units}$$

Step 2: Divide the increase by the January sales.

$$30 \div 120 = 0.25$$

Step 3: Convert to a percentage.

$$0.25 \times 100 = 25\%$$

The store's sales increased by 25% from January to February.

📄 🗨️ 🔄 ⋮

General practices when prompting>>

- Be very clear and specific
- Provide relevant context
- Give few examples
- Break complex tasks into steps
- Give the model a role when useful
- Iterate, refine, and verify important outputs

Building AI applications: RAGs and AI agents

Learn how to build AI applications for specific use cases

September 11, 2026

Analytics Solutions Bootcamp

45

From AI models to applications >>

A foundation model by itself has limitations:

- It may not know your organization's private or recent information.
- It cannot automatically access databases, APIs, or business systems.
- It may generate plausible but incorrect answers.
- It normally cannot take actions unless tools are made available.

Building AI applications: RAGs and AI agents

From AI models to applications >>

A practical AI application therefore combines:


Model + Context + Knowledge + Tools + Business Logic

Key takeaway:

An AI model *generates responses*. An AI application *connects the model to the **information and capabilities*** needed to solve a business problem.

Retrieval-Augmented Generation >>

Retrieval-Augmented Generation (RAG) is an approach that retrieves relevant information from an external knowledge source and provides it to a generative AI model as context before generating a response.




RAGs sample use cases





Some of the common use cases of RAGs are as follows.

- Company policy assistant
- Employee handbook Q&A
- Product documentation assistant
- Customer support knowledge base



Building AI applications: RAGs and AI agents

RAGs example



September 11, 2026

Analytics Solutions Bootcamp


50

RAGs example



We can keep adding new files that the model can reference for improved responses.

- Tip: Converting the files to markdown is less costly



Vector databases



When you have massive number of files that you want an AI model to use as reference, creating a vector database would be much more efficient.

Your raw files can be processed and converted to numerical vectors called embeddings. We use vector databases to store and retrieve these embeddings.

We won't cover the demo on creating vector databases.

AI Agents



An **AI agent** is an AI-powered system that can work toward a goal by deciding what steps to perform and which available tools to use. It carries out a task with three components: a **trigger**, a **process** that may include specialized skills, and **tools or systems** it can connect to.

Chatbot: Prompt → Response

AI Agent: Trigger → Defined Process → Use tools

AI Agents vs workflows



Traditional workflows	AI Agents
Steps are explicitly programmed	Steps and processes are provided via prompts
Fails on error	Agents can rectify errors
Tasks performed are fixed	Tasks performed may vary

Building AI applications: RAGs and AI agents

Example AI agents



Scheduled

Ask ChatGPT to schedule tasks, set reminders, or monitor for updates.

+ Schedule a task

🔒

Closing Portfolio Alert

Weekly · Next run in 21 hours

Midday Portfolio Alert

Weekly · Next run in 18 hours

Pre-Market Portfolio Brief

Weekly · Next run in 10 hours

Recommended

–

Daily brief

Personalized daily briefing with updates about the topics that I care about most

+

Email monitor

Scan my emails and let me know anything that needs my attention

+

Josh Dev PM progress

+

Closing Portfolio Alert

Weekly · Next run in 21 hours

Friday 2:58 PM

Market-close portfolio update — August 28

Your latest COL snapshot is more useful than the public quote feeds today, which are lagging the August 28 session. At roughly the close, your portfolio showed \$497,060.40 total account equity, \$21,207.22 buying power, and \$465,853.18 portfolio trade value. The portfolio was down about \$2,363.57 (-0.67%) on the day and \$13,854.89 (-2.91%) overall. Most importantly, your 100-share ABET purchase at \$36.70 executed, taking ABET to 1,000 shares and lowering its COL average cost from roughly \$32.71 to \$36.5233.

Today's market backdrop remains difficult. Yesterday's RSP rate hike to 5.00% triggered heavy de-risking, with the PSE closing Thursday at 6,004.58 (-2.16%) and net foreign selling reaching \$2.22 billion. Today the peso also opened beyond \$428 for the first time, reaching \$42.23 intraday early in the session. That combination—higher rates plus a weaker peso—remains a valuation headwind for equities, particularly REITs and other yield sensitive securities. @investor101

Based on your actual COL prices near today's close:

Headline	Price	Day/Session	Comment	Rating	My action
ABET	\$36.5233			completed, don't	

Scheduled

Ask ChatGPT to schedule tasks, set reminders, or monitor for updates.

+ Schedule a task

🔒

Closing Portfolio Alert

Weekly · Next run in 21 hours

Midday Portfolio Alert

Weekly · Next run in 18 hours

Pre-Market Portfolio Brief

Weekly · Next run in 10 hours

Recommended

–

Daily brief

Personalized daily briefing with updates about the topics that I care about most

+

Email monitor

Scan my emails and let me know anything that needs my attention

+

Josh Dev PM progress

+

Midday Portfolio Alert

Weekly · Next run in 18 hours

Friday 12:01 PM

Midday portfolio update — August 28

There is an important limitation today: the public market feeds I can verify are still showing August 27 as the latest trading session, with the PSE at 6,004.58 (-2.16%). One current market page explicitly reports the market as closed and labels August 27 as the last trading session. @101: Because I can't reliably exclude from August 28 noon quotes, I won't present yesterday's prices as today's live prices.

That said, valuations have moved enough that two of our accumulation targets are now very close.

Holding	Latest verified price	Rating	Outdated strategy view
ABET	\$37.10	BUY	First accumulation candidate
GLD	\$1,649	HOLD	Attractive yield, but directly bear exposure
LTC	\$15.16	HOLD	Wait closer to \$16
MBT	\$55.75	BUY	Now only \$5.75 above
NER	\$270.00	HOLD	Chopping, but cool

September 11, 2026

Analytics Solutions Bootcamp

55

Evaluating AI solutions

How do we know whether an AI solution is actually good enough to use?

Why evaluate AI solutions?



AI outputs are **probabilistic**. The same system can perform well on one task and poorly on another. Evaluation helps determine whether an AI solution is:

- Correct: Does it produce factually accurate results?
- Relevant: Does it address the user's request?
- Reliable: Does it perform consistently across different cases?
- Safe: Does it avoid harmful or inappropriate behavior?
- Useful: Does it actually help users accomplish the business objective?

What should we evaluate?



- **Model output quality:** accuracy, relevance, completeness, hallucinations
- **Application, system performance:** response latency, code safety, system costs
- **Business Outcomes:** time saved, increased productivity, customer satisfaction, resolution rate, reduced errors

Key takeaways on AI solutions >>

- A successful demo **does not necessarily mean a production-ready** AI solution.
- The AI must ultimately improve the **business process or outcome** it was designed for.
- AI evaluation is not a one-time test. It is a **continuous process** of measuring, improving, and monitoring the solution.

Responsible AI and AI governance

What can go wrong with AI, and what controls should organizations put around it?

How do we responsibly use AI >>

Responsible AI refers to the principles and practices used to design, develop, deploy, and use AI systems in ways that are trustworthy, safe, fair, and aligned with organizational and societal expectations.

- Key principles include: fairness, transparency, accountability, privacy, safety and reliability, and human oversight

Just because an AI solution can be built does not mean it should be deployed without appropriate controls.

Risks in AI solutions



AI Risk	Example
Hallucination	AI invents a company policy that does not exist
Bias and fairness	Hiring model systematically disadvantages certain applicants
Privacy	Employee personal information is exposed through an AI assistant

Risks in AI solutions



AI Risk	Example
Security	Prompt injection causes an agent to perform unintended actions
Lack of transparency	Users do not know that a decision was AI-generated
Over-reliance	Employee accepts an AI recommendation without verification

Risks in AI solutions



AI Risk	Example
Unauthorized actions	Agent performs actions beyond what the user intended
Intellectual property	Sensitive or copyrighted material is improperly used or generated

Responsible AI and AI governance

Risks in AI solutions



AI risks increases with **data access**, **capability**, and **autonomy**.

Primer to AI governance



AI governance is the framework of policies, roles, processes, standards, and controls used to manage how AI is developed, acquired, deployed, and monitored within an organization. Here are some examples of governance controls for AI.

- Approved AI tools/models, AI use policies
- Data access controls
- Risk assessment, testing and evaluation requirements
- Monitoring and incident management

Human in the loop



Human-in-the-loop (HITL) means incorporating human review, approval, or intervention into an AI-enabled process, particularly when decisions have significant consequences.

The appropriate level of human oversight should depend on the risk and impact of the AI decision or action.

Human in the loop's direction is focused on AI-first solutions

Responsible AI and AI governance

AI in the loop



Instead of inserting humans into machine workflows, insert machines into human ones.

AI in the loop is human first, then AI second.

In our processes identify the steps where AI can really help on without completely allocating the processes and the human roles to AI.

September 11, 2026

Analytics Solutions Bootcamp

68

<https://www.mindfoundry.ai/blog/ai-in-the-loop>

Responsible AI and AI governance

On responsible AI use



- **Start with the business problem, not with the technology.**
- **Build for value. Evaluate for performance. Govern for responsible use.**